

Semantic Epistemic Risk versus Confidence-Based Active Retrieval During Generation

Working Draft

August 2026

Abstract

Active retrieval during generation is established prior art: FLARE retrieves when forecast tokens have low model confidence [1]. We ask whether uncertainty is equivalent to epistemic risk. An initial 80-example controlled freeze finds that a transparent semantic trigger has 100% retrieval recall and 2.5% false activation across Qwen3-0.6B, SmolLM2-1.7B, and Gemma3-1B, while confidence is checkpoint-dependent. We then freeze a 192-example natural-language extension after rejecting four candidate freezes for lexical triviality, duplicates, source-key collisions, or inconsistent gold answers. On its 64-example test, the original shallow lexical accuracy is at most 85.94%. A matched tokenizer-aware audit changes that boundary: native Qwen, SmolLM2, and Gemma token-bigram TF-IDF classifiers each reach 100%, while a shared word/punctuation tokenizer reaches 87.5%. The unchanged semantic policy reaches 100% recall and 12.5% false activation; confidence recall is 96.88%, 81.25%, and 78.12% with false activation of 62.5%, 0%, and 9.38%. Advisory evidence leaves UCR of 18.75%, 12.5%, and 3.12%; an explicit support contract leaves 12.5%, 12.5%, and 0%. Deterministic runtime enforcement reaches zero UCR and full authorized coverage for every family. Prior rank-8 LoRAs have zero held-out recall, and the consumed natural freeze is not reused for adaptation. The result supports transparent runtime enforcement, but not superiority over stronger model-token lexical triggers. The tokenizer result exposes residual template structure and requires a harder independently frozen benchmark. A first public CRAG transfer audit confirms that requirement: on 320 questions paired with context-sufficient controls, the natural semantic rule has only 4.69% needed recall at zero false retrieval. A subsequent developmental Qwen3-0.6B slice evaluates 40 questions under matched external/context availability. Voluntary generic retrieval has zero external recall; confidence retrieval reaches 82.5% recall but retrieves on 45% of controls; and the unchanged semantic runtime has zero recall. Confidence automatically rescues 22/40 voluntary misses, while the registered oracle rescues 28/40. These public results falsify transfer of the present semantic trigger and bound, rather than establish, a deployable policy.

1 Question and Claim Boundary

The primary question is: *is epistemic risk distinct from model uncertainty?* Retrieval-required means that a reliable answer depends on an external or source-specific record rather than model weights or already supplied active context. This operational label is distinct from model confidence.

The secondary question is placement: *where should stable one-source interface knowledge live: context, weights, or runtime?* Phase B begins only after the Phase A source, benchmark, confidence probes, threshold objective, and evaluation are frozen. The paper does not claim active retrieval itself as novel and does not study multi-source routing.

2 Mechanism

$Q \rightarrow y_{1:t} \rightarrow \hat{y}_{t+1:t+k} \rightarrow \{T_{\text{conf}}, T_{\text{sem}}\} \rightarrow \text{retrieve} \rightarrow \text{evidence/status} \rightarrow \text{accept or abstain}.$

The FLARE-like controller greedily forecasts one short answer span, records the probability of every forecast token, triggers on the minimum probability, and regenerates with evidence. Canonical FLARE forecasts longer text, may mask uncertain spans, and iterates during long-form generation [1]; our single-opportunity controlled diagnostic is not a canonical reproduction.

The semantic controller uses answer-independent lexical, entity, temporal, freshness, source-dependence, and active-context rules. It has an explicit exception for “electric current” and treats compute requests as belonging to a different coprocessor. Rules never contain source answer values.

3 Controlled Source and Evidence Policy

The frozen source is `atlas-controlled-registry@2026.08.2-frozen-candidates`. Records have stable IDs, validity intervals, source version, and provenance. They cover stable, superseded familiar, unfamiliar private, unavailable, and conflicting facts. Requests contain entity, attribute, as-of date, forecast, source version, and resource bounds. Outcomes include SUPPORTED, CONTRADICTED, UNVERIFIED, STALE, AMBIGUOUS, and CONFLICT.

The runtime gate enforces support relative to this configured source; it does not establish truth. A unique supported value may be accepted. Missing, ambiguous, or conflicting evidence forces abstention rather than an unsupported factual commitment.

4 Frozen Measured-Quadrant Benchmark

We generate 256 candidates: 32 development and 224 candidate-test examples. Confidence is measured on pinned Qwen3-0.6B before test selection. The development threshold maximizes retrieval-requirement F1 minus 0.25 times retrieval rate and equals 0.84993. The candidate pool contains 25, 87, 22, and 90 examples in the four observed Qwen quadrants; the freezer selects exactly 20 from each. IDs are fixed before any condition evaluation.

Retrieval-required subclasses are fresh/current, source-version-specific, private, changed familiar, unavailable, conflicting, exact attribution, and structured one-source values. Retrieval-not-required subclasses are stable facts, supplied context, freshness distractors, quotations, hypotheticals, non-factual prose, and compute needs. The final benchmark contains 32 development and 80 untouched test examples.

Each model independently fits its threshold on the common development set: 0.84993 for Qwen, 0.14693 for SmolLM2, and 0.99971 for Gemma. Qwen and Gemma observe all four test quadrants. SmolLM2 places no test example below its fitted threshold, so its test set lacks low-confidence quadrants. This is a calibration result and limits cross-model quadrant claims.

5 Conditions and Metrics

Phase A has 13 conditions: LLM only, anti-hallucination prompting, upfront RAG, explicit retrieval, FLARE-like confidence, semantic, confidence OR semantic, confidence AND semantic, retrospective verification, evidence advisory, evidence plus abstention instruction, runtime epistemic gate, and oracle.

UCR is the fraction of retrieval-required opportunities that emit an unsupported factual commitment. Authorized Commitment Coverage is the fraction of non-abstention retrieval-required

Table 1: Untouched 80-example Phase A test. Acc.=exact accuracy, Ret.=retrieval rate, Cov.=Authorized Commitment Coverage.

Model	Condition	Acc.	Ret.	UCR	Cov.
Qwen	LLM only	.238	.000	1.000	.000
	FLARE-like	.625	.500	.675	.325
	Semantic	.613	.513	.250	.750
	Conf. OR semantic	.838	.763	.250	.750
	Runtime gate	.975	.763	.000	1.000
	Upfront RAG	.850	1.000	.250	.750
SmolLM2	LLM only	.463	.000	1.000	.000
	FLARE-like	.463	.000	1.000	.000
	Semantic	.575	.513	.775	.225
	Conf. OR semantic	.575	.513	.775	.225
	Runtime gate	.963	.513	.000	1.000
	Upfront RAG	.400	1.000	.775	.225
Gemma	LLM only	.463	.000	1.000	.000
	FLARE-like	.813	.900	.325	.675
	Semantic	.800	.513	.300	.700
	Conf. OR semantic	.825	.913	.300	.700
	Runtime gate	.775	.913	.000	1.000
	Upfront RAG	.825	1.000	.300	.700

opportunities answered with the authorized value. We also report exact accuracy, retrieval precision/recall and rate, confident-hallucination catch rate, quadrant retrieval, runtime enforcement, generated tokens, calls, and wall time. Accuracy intervals and complete threshold sweeps are machine-readable.

6 Phase A Results

All runs use greedy decoding, seed 17, pinned revisions, PyTorch 2.13.0+XPU, and Intel integrated graphics. Each family produces 1,456 prediction rows and 1,456 traces. Table 1 reports the primary conditions.

Confidence and semantic risk are not equivalent. Qwen confidence retrieves both low-confidence quadrants, giving 50% recall and 50% false intervention; the semantic rule retrieves all 40 required cases and one of 40 controls. SmolLM2’s fitted confidence retrieves nothing. Gemma confidence has 97.5% recall but 82.5% false activation. The same semantic policy has 100% recall and 2.5% false activation on all three because it is independent of checkpoint logits.

The runtime gate eliminates unsupported commitments in every family. It also shows why reliability and answer accuracy must be separated: Qwen and SmolLM2 improve sharply, whereas Gemma drops five accuracy points below upfront RAG. Enforcement prevents unsupported acceptance but cannot guarantee that the base answer or abstention policy maximizes exact-match utility.

7 Phase B: Context, Weights, or Runtime

The adapter dataset contains 160 balanced training rows and 40 development rows with disjoint entity namespaces and values. It excludes every held-out benchmark entity and answer. Targets contain either a four-field `retrieve` block or `NO_RETRIEVAL`; no target contains an answer value. The overlap audit has zero train/development, held-out entity, or held-out value collisions.




Figure 1: Retrieval cost against UCR and Authorized Commitment Coverage across three checkpoints. The runtime gate reaches zero UCR but has model-dependent accuracy consequences.

Table 2: Held-out 80-example placement evaluation. Sel.=selection accuracy, Int.=valid interface success, Rec.=retrieval recall, FA=false activation.

Model	Placement	Sel.	Int.	Rec.	FA	Prompt tok.	Gen. tok.
Qwen	Context	.500	.500	.000	.000	136.6	6.2
	Weights	.500	.500	.000	.000	67.6	6.0
	Runtime semantic	.988	.988	1.000	.025	101.4	10.2
SmolLM2	Context	.050	.013	.075	.975	172.2	31.4
	Weights	.500	.500	.000	.000	91.2	7.0
	Runtime semantic	.988	.988	1.000	.025	137.0	8.2
Gemma	Context	.863	.363	1.000	.275	139.0	24.4
	Weights	.500	.500	.000	.000	66.0	5.8
	Runtime semantic	.988	.988	1.000	.025	101.1	8.3

Rank-8 adapters train for two epochs on Qwen, SmolLM2, and Gemma. Final development losses are 0.00028, 0.0004, and 0.0001. This apparent fit does not transfer. Table 2 compares held-out policy selection. Interface success additionally requires the correct entity, attribute, date, and source. Prompt-token values are measured condition totals; the runtime semantic rows use the Phase A answer template and are not a matched estimate of policy-only prompt overhead.

All adapters collapse to NO_RETRIEVAL: zero recall, no false activation, and 50% balanced selection accuracy. Qwen context behaves the same. SmolLM2 context mostly over-triggers and often emits invalid blocks. Gemma context detects all needs but only 36.25% of all examples pass the complete interface criterion. Adding confidence to adapter outputs does not recover Qwen or SmolLM2 and inherits Gemma confidence’s 82.5% false activation.

Near-zero development loss alongside zero held-out recall is direct evidence of protocol overfitting to narrow synthetic phrasing. It falsifies the proposed weight-placement benefit under this training design. More diverse training may change that result, but adapting after inspecting these held-out failures would require a new freeze.

8 Natural-Language Robustness Iteration

We next generate 192 balanced examples: 96 training, 32 development, and 64 untouched test cases. Eight retrieval-required categories cover release-specific facts, private identifiers, authorita-




Figure 2: Interface success and false activation for context, weights, confidence, transparent runtime semantics, combinations, and oracle.

tive attribution, changed familiar facts, dynamic assignments, recorded fields, opaque fields, and post-reorganization assignments. Eight controls include supplied context, quoted freshness language, hypotheticals, electric current, historical-copy tasks, fictional sources, computation, and stable familiar facts. Train, development, and test use distinct base paraphrases. The freezer verifies zero exact duplicates, zero normalized template overlap, unique source entity–attribute keys, and consistent gold answers for split-specific controls.

Four pre-freeze diagnostics are retained but excluded from claims. Version 1 failed the lexical screen at 93.75%. Version 2 passed at 87.5% but contained repeated questions. Version 3 removed duplicates but exposed source-key collisions between two categories. Version 4 fixed source keys but exposed split-template/gold-answer mismatches in several controls. Version 5 fixes these defects. Bag-of-words, word n -gram TF–IDF, and character n -gram TF–IDF logistic regressions score 81.25%, 85.94%, and 81.25% on its test set. This passes the predeclared 90% triviality screen, but 85.94% remains high and bounds the robustness claim.

8.1 Tokenization and lexical trigger strength

We compare the current word and character representations with a small Unicode tokenizer that splits words and punctuation. Separate conditions use the exact pinned tokenizer of each checkpoint. Token methods receive unigram and unigram–bigram TF–IDF logistic classifiers. We also build BM25 exemplar indexes with unigram and bigram terms; $k \in \{1, 3, 5, 10\}$ and the retrieval threshold are selected on development only. Raw and normalized native pieces are retained separately. No LLM call or recurring context token is used. Per-example pieces, predictions, vocabulary and index sizes, unknown-piece rates, latency, confusion matrices, and subclass scores are machine-readable.

This falsifies a broad reading of semantic-rule superiority on both freezes. All three pre-specified model-aligned n -gram conditions are perfect, including on historical-copy controls. Natural-v5 BM25 is also competitive: normalized Qwen and SmolLM2 pieces reach 1.000, but native-piece BM25 falls to .725 with .550 false activation on the original freeze. Thus native tokenization can expose highly discriminative surface structure, while the BM25 transfer result shows that model pieces are not uniformly better. Eighteen natural-v5 and 26 original-freeze conditions tie at perfect development accuracy, so an arbitrary development tie winner is not evidence of a unique production choice. The proper conclusion is benchmark lexical leakage/headroom, not general semantic content in BPE tokens. Runtime enforcement still determines what happens after a trigger and still yields

Table 3: Matched trigger classification. Acc.=binary accuracy, Rec.=retrieval recall, FA=false activation. Native rows use raw model pieces and token unigrams+bigrams. The 80-example freeze uses its disjoint 160/40 policy train/development sets; natural-v5 uses its own 96/32 split.

Representation	Method	Original 80			Natural-v5		
		Acc.	Rec.	FA	Acc.	Rec.	FA
Current word	TF-IDF <i>n</i> -gram	.725	1.000	.550	.859	.969	.250
Current char	TF-IDF <i>n</i> -gram	.988	1.000	.025	.813	.875	.250
Shared NLP	TF-IDF <i>n</i> -gram	.725	1.000	.550	.875	1.000	.250
Qwen native	TF-IDF <i>n</i> -gram	1.000	1.000	.000	1.000	1.000	.000
SmolLM2 native	TF-IDF <i>n</i> -gram	1.000	1.000	.000	1.000	1.000	.000
Gemma native	TF-IDF <i>n</i> -gram	1.000	1.000	.000	1.000	1.000	.000
Semantic runtime	Fixed rules	.988	1.000	.025	.938	1.000	.125

Table 4: Natural-v5 test trigger and enforcement results. Rec.=retrieval recall, FA=false activation, Adv.=evidence advisory, Contract=explicit support contract.

Model	Conf. Rec.	Conf. FA	Sem. Rec.	Sem. FA	Adv. UCR	Contract UCR	Runtime UCR	Runtime Cov.
Qwen3-0.6B	.969	.625	1.000	.125	.188	.125	.000	1.000
SmolLM2-1.7B	.813	.000	1.000	.125	.125	.125	.000	1.000
Gemma3-1B	.781	.094	1.000	.125	.031	.000	.000	1.000

zero UCR when recall is one.

Figure 3 shows that confidence calibration changes again: Qwen nearly reaches full recall but over-triggers, whereas SmolLM2 and Gemma miss required cases. The frozen semantic policy catches every required case but falsely retrieves four historical-copy controls containing both private IDs and reorganization language. We do not tune this error after inspecting test.

Evidence use is model-dependent. Advisory evidence yields authorized coverage of 81.25%, 87.5%, and 96.88%. The explicit support contract lowers UCR but Qwen over-abstains, reducing its authorized coverage to 21.88%; SmolLM2 and Gemma retain 87.5% and 81.25%. Runtime enforcement copies a unique supported value or abstains otherwise, reaching zero UCR and full authorized coverage for all three. Its exact accuracy is 81.25%, 90.62%, and 68.75%, versus 37.5%, 46.88%, and 25% without retrieval and 87.5%, 96.88%, and 75% for the source-copy oracle. Retrospective verification detects all wrong required forecasts but, without a forced rewrite, leaves UCR at 100%. Detection is therefore not correction. Per-condition token, call, retrieval-time, and wall-time totals remain machine-readable in the released summaries and traces.

A deterministic four-document controller stress test has three factual opportunities per document: a control, a later emerging evidence need, and reuse of the same evidence. It records 100% early prospective catch, 100% late retrospective detection, 50% reuse across required opportunities, no repeated unnecessary retrieval, and zero runtime UCR. This tests interruption and cache mechanics, not free-running long-form generation; it is not evidence that models discover all factual opportunities in unconstrained prose.

The learned-policy gate remains NO-GO. Natural-v5 is now consumed, shallow lexical accuracy remains high, semantic rules leave only false-activation headroom, and the prior adapters failed held-out retrieval despite near-zero development loss. Reopening LoRA requires richer training data and a new test freeze created before adaptation.




Figure 3: Checkpoint confidence versus the unchanged transparent semantic policy on the 64-example natural-v5 test.




Figure 4: Natural-v5 enforcement ladder and deterministic multi-opportunity controller stress test. Runtime enforcement is reliable relative to the configured source; it is not a truth oracle.

9 Public CRAG Transfer Audit

We freeze 320 questions from a pinned 2,700-row CRAG index, stratified jointly over domain, question type, static/dynamic class, answerability, and source split. The subset contains all five domains, all eight question types, 116 static, 102 slow-changing, 79 fast-changing, and 23 real-time questions. Each question yields two matched rows: an external condition labeled retrieval required and a control that supplies the gold answer in active context. This tests the operational equation “evidence requirement minus available evidence equals retrieval action” without learning source choice. Artifacts retain IDs, strata, source rows, and content hashes but do not redistribute questions or answers.

The natural rule correctly suppresses every matched context control, but it recognizes only 15 external questions. The legacy rule also reacts to temporal content injected with some supplied answers, so its control false-retrieval rate exceeds its external recall. Treating every non-static row as risky improves recall to 0.638 and perfectly suppresses matched controls, yet misses every static source-backed question. Therefore semantic policy behavior from the synthetic freeze does not

Table 5: CPU-only CRAG trigger transfer on 320 external questions and 320 matched context controls. Recall is on external rows; FA is false retrieval on controls; UCR is unsupported commitment risk under trigger enforcement.

Condition	Recall	FA	Accuracy	UCR	Retrieval rate
No retrieval	.000	.000	.500	1.000	.000
Upfront retrieval	1.000	1.000	.500	.000	1.000
Generic <code>__retrieve</code> , oracle timed	1.000	.000	1.000	.000	.500
Legacy semantic rule	.141	.206	.467	.859	.173
Natural semantic rule	.047	.000	.523	.953	.023
Dynamic metadata proxy	.638	.000	.819	.363	.319
Registered oracle	1.000	.000	1.000	.000	.500




Figure 5: The synthetic semantic rules do not transfer to CRAG. The metadata proxy is diagnostic, and the registered oracle is not deployable.

transfer, and “dynamic” is not equivalent to “needs evidence.”

The all/none conditions expose the paired benchmark’s operating bounds. Upfront retrieval attains full needed-retrieval recall only by retrieving on every context-sufficient control. Passing the 320 registered external decisions through the shared four-tool gateway’s `__retrieve` interface produces 320 accepted calls and agrees with the oracle on every row. That condition is deliberately oracle timed: it validates a stable tool name and runtime handoff, not a model’s ability to request retrieval.

This 320-question stage remains a trigger and transport audit rather than a RAG result. The next developmental stage therefore freezes a disjoint 20-question calibration slice and 40-question evaluation slice, balanced across CRAG’s four static/dynamic classes. Each evaluation question is run both without evidence and with its reference answer in active context. The bounded retrieval backend returns the frozen CRAG reference answer; consequently this tests trigger, handoff, evidence use, and answer generation, but not open-web search quality or multi-source routing.

The voluntary generic-tool prompt never retrieves on the 40 external rows. Its five-percent false-retrieval rate comes only from controls, so voluntary tool availability is not a retrieval policy. Confidence is substantially more sensitive, retrieving on 33/40 external rows, but also on 18/40 controls. Relative to all 40 voluntary external misses, confidence retrieves and answers 22 correctly,

Table 6: Model-facing Qwen3-0.6B CRAG development slice (40 questions, 80 matched rows). Ext./Ctx. are answer accuracy under external and supplied-context availability; recall is on external rows and FA on controls.

Condition	Ext.	Ctx.	Recall	FA	UCR
LLM only	.025	.025	.000	.000	1.000
Upfront retrieval	.700	.700	1.000	1.000	.000
Voluntary __retrieve	.050	.075	.000	.050	.800
Semantic runtime	.025	.025	.000	.000	1.000
Confidence trigger	.550	.275	.825	.450	.175
Generic retrieval block	.050	.200	.025	.400	.950
Registered oracle	.700	.025	1.000	.000	.000
Runtime answer copy	1.000	1.000	1.000	.000	.000




Figure 6: Answer accuracy and retrieval-needed recall on the matched CRAG model slice. Runtime copy is a deterministic upper-bound control, not model behavior.

for 55% Automatic Rescue Rate. The registered oracle retrieves all misses and answers 28 correctly, giving a 70% upper bound for this model and grounded prompt. Thus seven oracle-routed answers still fail generation; retrieval correctness does not imply answer correctness. Runtime answer copy is reported only to verify evidence transport and scoring.

Most importantly, the unchanged semantic policy triggers on none of these matched model rows and rescues no voluntary misses. This is consistent with the low recall in the larger CPU transfer audit and directly falsifies transfer of the synthetic semantic trigger. The result supports runtime enforcement once a valid evidence need is identified, but it does not show that the current semantic detector identifies that need on CRAG.

10 Paper 2.5 Gate

The pre-registered gate asks whether at least two families reduce UCR with the runtime gate while retrieving less often than upfront RAG. Qwen, SmolLM2, and Gemma all pass; Paper 2.5 is therefore GO. This authorizes testing source heterogeneity. It does not authorize learned source

routing, because Phase B provides no evidence that the retrieval policy should move into weights.

11 Falsification and Limitations

The primary result would be weakened if a better-calibrated FLARE implementation matches semantic routing at equal cost, if natural-language semantic rules lose precision, or if the confident-risk quadrants disappear on new data. The present result instead finds checkpoint-dependent confidence behavior and stable rule recall, but only on one controlled source and mostly synthetic short answers.

The benchmark is selected using Qwen confidence, so its exact 20-per-quadrant balance does not transfer to other checkpoints. SmolLM2 lacks low-confidence test examples. There is one seed per model and no ranking noise. The natural extension passed its limited lexical screen only narrowly; the tokenizer-aware audit subsequently finds perfect model-token classifiers, so that screen was insufficient. Its long-form component is a deterministic controller simulation, not unconstrained generation. The semantic cue distribution remains controlled. The runtime gate is deterministic and answer-aware relative to source evidence, not a production truth system. Full-prefix generation latency is not production latency.

Phase B uses narrow train templates despite strict entity/value leakage controls; its failure may be covariate shift rather than a general limitation of LoRA. Conversely, development loss is clearly not evidence of interface generalization. The ICL baseline uses textual examples rather than native tool schemas. No condition trains or tests multi-source selection.

The 320-question CRAG checkpoint uses index metadata and matched gold-answer context controls, not retrieved passages. Its dynamic proxy relies on benchmark metadata that a deployment may not possess. The model-facing slice uses one Qwen checkpoint, 40 development questions, one deterministic decode seed, and reference answers as bounded retrieval evidence. Confidence is calibrated on a separate 20-question slice, but no untouched confirmatory CRAG test has yet been run. Exact-match scoring can undercount semantically equivalent answers. The zero semantic recall and high confidence false-retrieval rate are diagnostic failures to fix, not population estimates or a ranking of complete retrieval systems.

12 Reproduction

Paper-specific code is under `src/ccpu/paper1_5`. From the repository root, activate the validated XPU environment and run:

```
$py = 'C:\Users\j.simao\.venvs\modal-llm-xpu\Scripts\python.exe'
$env:PYTHONPATH = (Resolve-Path src).Path
& $py -m ccpu paper1.5 freeze-next \
  --config configs/paper1_5/next_iter_freeze_qwen_xpu.json \
  --output-dir artifacts/paper1_5/next_iter/freeze_qwen_v2
& $py -m ccpu paper1.5 run-hf \
  --config configs/paper1_5/next_iter_qwen_xpu.json \
  --output-dir artifacts/paper1_5/next_iter/qwen_v2
python -m ccpu paper1.5 analyze-next \
  --config configs/paper1_5/next_iter_analysis.json \
  --output-dir artifacts/paper1_5/next_iter/analysis_v1
& $py -m ccpu paper1.5 train-policy-lora <arguments in README>
```

```

python -m ccpu paper1.5 analyze-policy \
  --config configs/paper1_5/policy_placement.json \
  --output-dir artifacts/paper1_5/next_iter/policy_analysis_v1
python -m ccpu paper1.5 freeze-natural \
  --config configs/paper1_5/natural_robustness_xpu.json \
  --output-dir artifacts/paper1_5/natural_v5/data
& $py -m ccpu paper1.5 run-natural <natural-v5 paths and model label>
python -m ccpu paper1.5 analyze-natural \
  --config configs/paper1_5/natural_analysis.json \
  --output-dir artifacts/paper1_5/natural_v5/analysis
& $py -m ccpu paper1.5 compare-natural-tokenizers \
  --benchmark artifacts/paper1_5/natural_v5/data/benchmark.jsonl \
  --tokenizer-config configs/common/tokenizer_triggers.json \
  --output-dir artifacts/paper1_5/tokenizer_triggers/natural_v5
& $py -m ccpu paper1.5 compare-natural-tokenizers \
  --benchmark artifacts/paper1_5/next_iter/qwen_v2/benchmark.jsonl \
  --train artifacts/paper1_5/next_iter/policy_data_v2/train.jsonl \
  --dev artifacts/paper1_5/next_iter/policy_data_v2/dev.jsonl \
  --tokenizer-config configs/common/tokenizer_triggers.json \
  --output-dir artifacts/paper1_5/tokenizer_triggers/original80
& $py -m ccpu paper1.5 freeze-public-crag \
  --config configs/paper1_5/public_crag.json \
  --cache-root C:\path\to\pinned-hf-cache \
  --output-dir artifacts/paper1_5/public_benchmarks/crag_v1/data
& $py -m ccpu paper1.5 analyze-public-crag \
  --config configs/paper1_5/public_crag.json \
  --cache-root C:\path\to\pinned-hf-cache \
  --selection artifacts/paper1_5/public_benchmarks/crag_v1/data/selection.jsonl \
  --output-dir artifacts/paper1_5/public_benchmarks/crag_v1/analysis
& $py -m ccpu paper1.5 freeze-public-crag-model \
  --source-selection artifacts/paper1_5/public_benchmarks/crag_v1/data/selection.jsonl \
  --output-dir artifacts/paper1_5/public_benchmarks/crag_model_v1/data
& $py -m ccpu paper1.5 run-public-crag-model \
  --config configs/paper1_5/public_crag_model_qwen_xpu.json \
  --cache-root C:\path\to\pinned-hf-cache \
  --selection artifacts/paper1_5/public_benchmarks/crag_model_v1/data/selection.jsonl \
  --output-dir artifacts/paper1_5/public_benchmarks/crag_model_v1/qwen

```

Manifests hash configurations, benchmark/source snapshots, predictions, traces, summaries, and adapter files and record checkpoint revisions, device, package, repository, and environment provenance.

13 Conclusion

Semantic epistemic risk adds information beyond checkpoint uncertainty in the controlled one-source regime and retains full recall under a more varied natural-language freeze, but that behavior does not transfer to CRAG. It does not beat model-aligned token n -gram triggers: their perfect scores reveal that the present freezes remain lexically separable under stronger representations.

Advisory and instructed evidence use remain checkpoint-dependent; deterministic runtime enforcement eliminates unsupported commitments for all three small families. The placement extension is negative: context is inconsistent and the tested adapters overfit. Exact source mechanics and enforcement belong in runtime under these freezes, and trigger selection should remain CPU-side and auditable until a richer independently frozen benchmark separates lexical cues from epistemic semantics. On the developmental CRAG slice, confidence provides nonzero rescue but at high control retrieval cost; the current semantic and voluntary policies fail. The next claim-bearing step is therefore a new policy and untouched freeze, not a retrospective reinterpretation of this negative result.

References

- [1] Z. Jiang et al. Active retrieval augmented generation. *EMNLP*, 2023. <https://arxiv.org/abs/2305.06983>.